

收藏 | 深度学习优化算法：从SGD到AdamW原理和代码解读

计算机视觉联盟 Yesterday



计算机视觉联盟

团队成员为国内各大高校博士，专注于深度学习、机器学习、图像解译、人工...

115篇原创内容

Official Account

点上方 [计算机视觉联盟](#) 获取更多干货

仅作学术分享，不代表本公众号立场，侵权联系删除

转载于：作者 | 科技猛兽

编辑 | 极市平台

AI博士笔记系列推荐

[周志华《机器学习》手推笔记正式开源！可打印版本附pdf下载链接](#)

本文思想来自下面这篇大佬的文章：

Juliuszh：一个框架看懂优化算法之异同 SGD/AdaGrad/Adam

<https://zhuanlan.zhihu.com/p/32230623>

主要是对深度学习各种优化器 (从SGD到AdamW) 使用统一的框架做一次整理，本文相比于链接从源代码的角度理解这些优化器的思路。

代码来自 PyTorch1.7.0 官方教程：

<https://pytorch.org/docs/1.7.0/optim.html>

首先我们来回顾一下各类优化算法。

深度学习优化算法经历了 SGD -> SGDM -> NAG -> AdaGrad -> AdaDelta -> Adam -> Nadam -> AdamW 这样的发展历程。Google一下就可以看到很多的教程文章，详细告诉你这些算法是如何一步一步演变而来的。在这里，我们换一个思路，用一个框架来梳理所有的优化算法，做一个更加高屋建瓴的对比。

- 统一框架：

首先定义：待优化参数： w ，目标函数： $f(w)$ ，初始学习率 α 。

而后，开始进行迭代优化。在每个epoch t ：

1 计算目标函数关于当前参数的梯度：

$$g_t = \nabla f(w_t) \quad (1)$$

2 根据历史梯度计算一阶动量和二阶动量：

$$m_t = \phi(g_1, g_2, \dots, g_t); V_t = \psi(g_1, g_2, \dots, g_t) \quad (2)$$

3 计算当前时刻的下降梯度：

$$\eta_t = \alpha \cdot m_t / \sqrt{V_t} \quad (3)$$

4 根据下降梯度进行更新：

$$w_{t+1} = w_t - \eta_t \quad (4)$$

掌握了这个框架，你可以轻轻松松设计自己的优化算法。

我们拿着这个框架，来照一照各种玄乎其玄的优化算法的真身。步骤3, 4对于各个算法都是一致的，主要的差别就体现在1和2上，也就是计算一阶动量 m_t 和二阶动量 V_t 时采用不同的套路。当计算好二者之后，都是使用固定的学习率 α 与二者作用得到当前时刻的下降梯度 η_t ，进而最后更新参数。

在所有优化器的代码里面有一些函数的作用是相通的：

共性的方法有：

- `add_param_group` (param_group)：把参数放进优化器中，这在 Fine-tune 预训练网络时很有用，因为可以使冻结层可训练并随着训练的进行添加到优化器中。
- `load_state_dict` (state_dict)：把优化器的状态加载进去。
- `state_dict` ()：返回优化器的状态，以dict的形式返回。
- `step` (closure=None)：优化一步参数。
- `zero_grad` (set_to_none=False)：把所有的梯度值设为0。

使用方法：

```
for input, target in dataset:
    def closure():
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
        return loss
    optimizer.step(closure)
```

下面正式开始。

SGD

先来看SGD。SGD没有动量的概念，也就是说：

$$m_t = g_t; V_t = I^2 \quad (5)$$

代入步骤3，可以看到下降梯度就是最简单的

$$\eta_t = \alpha \cdot g_t \quad (6)$$

SGD最大的缺点是下降速度慢，而且可能会在沟壑的两边持续震荡，停留在一个局部最优

点。

SGD with Momentum

为了抑制SGD的震荡，SGDM认为梯度下降过程可以加入惯性。下坡的时候，如果发现是陡坡，那就利用惯性跑的快一些。SGDM全称是SGD with momentum，在SGD基础上引入了一阶动量：

$$m_t = \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t \quad (7)$$

一阶动量是各个时刻梯度方向的指数移动平均值，约等于最近 $1/(1 - \beta_1)$ 个时刻的梯度向量值的平均值。

也就是说， t 时刻的下降方向，不仅由当前点的梯度方向决定，而且由此前累积的下降方向决定。 β_1 的经验值为0.9，这就意味着下降方向主要是此前累积的下降方向，并略微偏向当前时刻的下降方向。想象高速公路上汽车转弯，在高速向前的同时略微偏向，急转弯可是要出事的。

SGD with Nesterov Acceleration

SGD 还有一个问题是困在局部最优的沟壑里面震荡。想象一下你走到一个盆地，四周都是略高的小山，你觉得没有下坡的方向，那就只能待在这里了。可是如果你爬上高地，就会发现外面的世界还很广阔。因此，我们不能停留在当前位置去观察未来的方向，而要向前一步、多看一步、看远一些。

NAG全称Nesterov Accelerated Gradient，是在SGD、SGD-M的基础上的进一步改进，改进点在于步骤1。我们知道在时刻 t 的主要下降方向是由累积动量决定的，自己的梯度方向说了也不算，那与其看当前梯度方向，不如先看看如果跟着累积动量走了一步，那个时候再怎么走。因此，NAG在步骤1，不计算当前位置的梯度方向，而是计算如果按照累积动量走了一步，那个时候的下降方向：

$$g_t = \nabla f(w_t - \beta_1 \cdot m_{t-1} / \sqrt{V_{t-1}}) \quad (8)$$

然后用下一个点的梯度方向，与历史累积动量相结合，计算步骤2中当前时刻的累积动量。

定义优化器：

```
CLASS torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0, weig
```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate，相当于是统一框架中的 α 。
- **momentum** (float, optional) – 动量参数。(默认值：0)
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty)
(默认值：0)
- **dampening** (float, optional) – dampening for momentum (默认值：0)
- **nesterov** (bool, optional) – 允许 Nesterov momentum (默认值：False)

FLOAT：<https://docs.python.org/3/library/functions.html#float>

bool:<https://docs.python.org/3/library/functions.html#bool>

源码解读：

```
import torch
from .optimizer import Optimizer, required
```

```
[docs]class SGD(Optimizer):
    r"""Implements stochastic gradient descent (optionally with momentum).

    Nesterov momentum is based on the formula from
    `On the importance of initialization and momentum in deep learning`__.

    Args:
        params (iterable): iterable of parameters to optimize or dicts defining
            parameter groups
        lr (float): learning rate
        momentum (float, optional): momentum factor (default: 0)
        weight_decay (float, optional): weight decay (L2 penalty) (default: 0)
        dampening (float, optional): dampening for momentum (default: 0)
        nesterov (bool, optional): enables Nesterov momentum (default: False)
```

Example:

```
>>> optimizer = torch.optim.SGD(model.parameters(), lr=0.1, momentum=0.9)
>>> optimizer.zero_grad()
>>> loss_fn(model(input), target).backward()
>>> optimizer.step()
```

__ <http://www.cs.toronto.edu/%7Ehinton/absps/momentum.pdf>

.. note::

The implementation of SGD with Momentum/Nesterov subtly differs from Sutskever et. al. and implementations in some other frameworks.

Considering the specific case of Momentum, the update can be written as

```
.. math::
    \begin{aligned}
        v_{t+1} &= \mu * v_t + g_{t+1}, \\
        p_{t+1} &= p_t - \text{lr} * v_{t+1},
    \end{aligned}
```

where p , g , v and μ denote the parameters, gradient, velocity, and momentum respectively.

This is in contrast to Sutskever et. al. and other frameworks which employ an update of the form

```
.. math::
    \begin{aligned}
```

$$\begin{aligned} v_{t+1} &= \mu * v_t + \text{lr} * g_{t+1}, \\ p_{t+1} &= p_t - v_{t+1}. \end{aligned}$$

The Nesterov version is analogously modified.

```

"""

def __init__(self, params, lr=required, momentum=0, dampening=0,
              weight_decay=0, nesterov=False):
    if lr is not required and lr < 0.0:
        raise ValueError("Invalid learning rate: {}".format(lr))
    if momentum < 0.0:
        raise ValueError("Invalid momentum value: {}".format(momentum))
    if weight_decay < 0.0:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))

    defaults = dict(lr=lr, momentum=momentum, dampening=dampening,
                    weight_decay=weight_decay, nesterov=nesterov)
    if nesterov and (momentum <= 0 or dampening != 0):
        raise ValueError("Nesterov momentum requires a momentum and zero dampening")
    super(SGD, self).__init__(params, defaults)

def __setstate__(self, state):
    super(SGD, self).__setstate__(state)
    for group in self.param_groups:
        group.setdefault('nesterov', False)

[docs] @torch.no_grad()
def step(self, closure=None):
    """Performs a single optimization step.

    Arguments:
        closure (callable, optional): A closure that reevaluates the model
            and returns the loss.

    """
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        weight_decay = group['weight_decay']
        momentum = group['momentum']
        dampening = group['dampening']
        nesterov = group['nesterov']

        for p in group['params']:
            if p.grad is None:
                continue
            d_p = p.grad
            if weight_decay != 0:
                d_p = d_p.add(p, alpha=weight_decay)
            if momentum != 0:
                param_state = self.state[p]
                if 'momentum_buffer' not in param_state:

```

```

        buf = param_state['momentum_buffer'] = torch.clone(d_p).data
    else:
        buf = param_state['momentum_buffer']
        buf.mul_(momentum).add_(d_p, alpha=1 - dampening)
    if nesterov:
        d_p = d_p.add(buf, alpha=momentum)
    else:
        d_p = buf

    p.add_(d_p, alpha=-group['lr'])

return loss

```

这里通过 $d_p = p.grad$ 得到每个参数的梯度，也就是1式的 g_t 。

如果使用 `weight_decay` 的话，那么相当于目标函数加上 $\frac{1}{2} \lambda ||W||^2$ ，所以相当于是梯度相当于要再加上 λW ，所以使用了 $d_p = d_p.add(p, \text{alpha}=\text{weight_decay})$ 。通过 $\text{buf.mul_}(\text{momentum}).\text{add_}(d_p, \text{alpha}=1 - \text{dampening})$ 来计算动量，momentum参数 β_1 一般取0.9，就相当于之前的动量buf乘以 $\beta_1 = 0.9$ ，再加上此次的梯度 d_p 乘以 $(1 - \beta_1) = 0.1$ 。

如果不通过nesterov方式更新参数，那么3式中的 η_t 就相当于上一步计算出的动量 m_t 了。如果通过nesterov方式更新参数，那么3式中的 η_t 就相当于 $g_t + m_t * \beta_1$ ，和不用nesterov方式相比，相差了。

最后通过 $p.add_(d_p, \text{alpha}=-\text{group}['lr'])$ 更新梯度，相当于是上面的3式。

AdaGrad

此前我们都没有用到二阶动量。二阶动量的出现，才意味着“自适应学习率”优化算法时代的到来。SGD及其变种以同样的学习率更新每个参数，但神经网络往往包含大量的参数，这些参数并不是总会用得到（想想大规模的embedding）。对于经常更新的参数，我们已经积累了大量关于它的知识，不希望被单个样本影响太大，希望学习速率慢一些；对于偶尔更新的参数，我们了解的信息太少，希望能从每个偶然出现的样本身上多学一些，即学习速率大一些。

怎么样去度量历史更新频率呢？那就是二阶动量——该维度上，迄今为止所有梯度值的平方和：

$$V_t = \sum_{\tau=1}^t g_{\tau}^2 \quad (9)$$

我们再回顾一下步骤3中的下降梯度：

$$\eta_t = \alpha \cdot m_t / \sqrt{V_t} \quad (3)$$

可以看出，此时实质上的学习率由 α 变成了 $\alpha/\sqrt{V_t}$ 。一般为了避免分母为0，会在分母上加一个小的平滑项。因此 $\sqrt{V_t}$ 是恒大于0的，而且参数更新越频繁，二阶动量越大，学习率就越小。

这一方法在稀疏数据场景下表现非常好。但也存在一些问题：因为 $\sqrt{V_t}$ 是单调递增的，会使得学习率单调递减至0，可能会使得训练过程提前结束，即便后续还有数据也无法学到必要的知识。

定义优化器：

```
CLASS torch.optim.Adagrad(params,lr=0.01,lr_decay=0,weight_decay=0,initial_accumulat
```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate – 相当于是统一框架中的 α 。
- **lr_decay**(float,optional) – 学习率衰减 (默认值：0)
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty) (默认值：0)
- **eps**(float,optional)：防止分母为0的一个小数 (默认值：1e-10)

源码解读：

```
[docs]class Adagrad(Optimizer):
    """Implements Adagrad algorithm.

    It has been proposed in `Adaptive Subgradient Methods for Online Learning
    and Stochastic Optimization`_.

    Arguments:
        params (iterable): iterable of parameters to optimize or dicts defining
            parameter groups
        lr (float, optional): learning rate (default: 1e-2)
        lr_decay (float, optional): learning rate decay (default: 0)
        weight_decay (float, optional): weight decay (L2 penalty) (default: 0)
        eps (float, optional): term added to the denominator to improve
```



```

        numerical stability (default: 1e-10)

.. _Adaptive Subgradient Methods for Online Learning and Stochastic
   Optimization: http://jmlr.org/papers/v12/duchilla.html
"""

def __init__(self, params, lr=1e-2, lr_decay=0, weight_decay=0, initial_accumula
    if not 0.0 <= lr:
        raise ValueError("Invalid learning rate: {}".format(lr))
    if not 0.0 <= lr_decay:
        raise ValueError("Invalid lr_decay value: {}".format(lr_decay))
    if not 0.0 <= weight_decay:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))
    if not 0.0 <= initial_accumulator_value:
        raise ValueError("Invalid initial_accumulator_value value: {}".format(in
    if not 0.0 <= eps:
        raise ValueError("Invalid epsilon value: {}".format(eps))

    defaults = dict(lr=lr, lr_decay=lr_decay, eps=eps, weight_decay=weight_decay
                    initial_accumulator_value=initial_accumulator_value)
    super(Adagrad, self).__init__(params, defaults)

    for group in self.param_groups:
        for p in group['params']:
            state = self.state[p]
            state['step'] = 0
            state['sum'] = torch.full_like(p, initial_accumulator_value, memory_

def share_memory(self):
    for group in self.param_groups:
        for p in group['params']:
            state = self.state[p]
            state['sum'].share_memory_()

[docs] @torch.no_grad()
def step(self, closure=None):
    """Performs a single optimization step.

    Arguments:
        closure (callable, optional): A closure that reevaluates the model
            and returns the loss.
    """
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        params_with_grad = []
        grads = []
        state_sums = []
        state_steps = []

        for p in group['params']:
            if p.grad is not None:

```

```

        params_with_grad.append(p)
        grads.append(p.grad)
        state = self.state[p]
        state_sums.append(state['sum'])
        # update the steps for each param group update
        state['step'] += 1
        # record the step after step update
        state_steps.append(state['step'])

    F.adagrad(params_with_grad,
              grads,
              state_sums,
              state_steps,
              group['lr'],
              group['weight_decay'],
              group['lr_decay'],
              group['eps'])

    return loss

```

AdaDelta / RMSProp

由于AdaGrad单调递减的学习率变化过于激进，我们考虑一个改变二阶动量计算方法的策略：不累积全部历史梯度，而只关注过去一段时间窗口的下降梯度。这也就是AdaDelta名称中Delta的来历。

修改的思路很简单。前面我们讲到，指数移动平均值大约就是过去一段时间的平均值，因此我们用这一方法来计算二阶累积动量：

$$V_t = \beta_2 * V_{t-1} + (1 - \beta_2)g_t^2 \quad (10)$$

接下来还是步骤3：

$$\eta_t = \alpha \cdot g_t / \sqrt{V_t} \quad (11)$$

这就避免了二阶动量持续累积、导致训练过程提前结束的问题了。

RMSProp

定义优化器：

```

CLASS torch.optim.RMSprop(params, lr=0.01, alpha=0.99, eps=1e-08, weight_decay=0, mo

```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate – 相当于是统一框架中的 α 。
- **momentum** (float, optional) – 动量参数。(默认值：0)。
- **alpha**(float, optional) – 平滑常数 (默认值：0.99)。
- **centered**(bool, optional) – if **True** , compute the centered RMSProp, the gradient is normalized by an estimation of its variance, 就是这一项是 True 的话就把方差使用梯度作归一化。
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty) (默认值：0)
- **eps**(float, optional)：防止分母为0的一个小数 (默认值：1e-10)

源码解读：

```
import torch
from .optimizer import Optimizer
```

```
[docs]class RMSprop(Optimizer):
    r"""Implements RMSprop algorithm.
```

```
    Proposed by G. Hinton in his
    `course <https://www.cs.toronto.edu/~tijmen/csc321/slides/lecture\_slides\_lec6.pdf>
```

```
    The centered version first appears in `Generating Sequences
    With Recurrent Neural Networks <https://arxiv.org/pdf/1308.0850v5.pdf>`.
```

```
    The implementation here takes the square root of the gradient average before
    adding epsilon (note that TensorFlow interchanges these two operations). The eff
    learning rate is thus  $\alpha / (\sqrt{v} + \epsilon)$  where  $\alpha$ 
    is the scheduled learning rate and  $v$  is the weighted moving average
    of the squared gradient.
```

```
    Arguments:
```

```
        params (iterable): iterable of parameters to optimize or dicts defining
            parameter groups
        lr (float, optional): learning rate (default: 1e-2)
        momentum (float, optional): momentum factor (default: 0)
        alpha (float, optional): smoothing constant (default: 0.99)
```

```

    eps (float, optional): term added to the denominator to improve
        numerical stability (default: 1e-8)
    centered (bool, optional) : if ``True``, compute the centered RMSProp,
        the gradient is normalized by an estimation of its variance
    weight_decay (float, optional): weight decay (L2 penalty) (default: 0)

"""

def __init__(self, params, lr=1e-2, alpha=0.99, eps=1e-8, weight_decay=0, momentum
    if not 0.0 <= lr:
        raise ValueError("Invalid learning rate: {}".format(lr))
    if not 0.0 <= eps:
        raise ValueError("Invalid epsilon value: {}".format(eps))
    if not 0.0 <= momentum:
        raise ValueError("Invalid momentum value: {}".format(momentum))
    if not 0.0 <= weight_decay:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))
    if not 0.0 <= alpha:
        raise ValueError("Invalid alpha value: {}".format(alpha))

    defaults = dict(lr=lr, momentum=momentum, alpha=alpha, eps=eps, centered=cent
    super(RMSprop, self).__init__(params, defaults)

def __setstate__(self, state):
    super(RMSprop, self).__setstate__(state)
    for group in self.param_groups:
        group.setdefault('momentum', 0)
        group.setdefault('centered', False)

[docs] @torch.no_grad()
def step(self, closure=None):
    """Performs a single optimization step.

    Arguments:
        closure (callable, optional): A closure that reevaluates the model
            and returns the loss.

    """
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:
                continue
            grad = p.grad
            if grad.is_sparse:
                raise RuntimeError('RMSprop does not support sparse gradients')
            state = self.state[p]

            # State initialization
            if len(state) == 0:
                state['step'] = 0
                state['square_avg'] = torch.zeros_like(p, memory_format=torch.pr

```

```

        if group['momentum'] > 0:
            state['momentum_buffer'] = torch.zeros_like(p, memory_format
        if group['centered']:
            state['grad_avg'] = torch.zeros_like(p, memory_format=torch.

square_avg = state['square_avg']
alpha = group['alpha']

state['step'] += 1

if group['weight_decay'] != 0:
    grad = grad.add(p, alpha=group['weight_decay'])

square_avg.mul_(alpha).addcmul_(grad, grad, value=1 - alpha)

if group['centered']:
    grad_avg = state['grad_avg']
    grad_avg.mul_(alpha).add_(grad, alpha=1 - alpha)
    avg = square_avg.addcmul(grad_avg, grad_avg, value=-1).sqrt().a
else:
    avg = square_avg.sqrt().add_(group['eps'])

if group['momentum'] > 0:
    buf = state['momentum_buffer']
    buf.mul_(group['momentum']).addcdiv_(grad, avg)
    p.add_(buf, alpha=-group['lr'])
else:
    p.addcdiv_(grad, avg, value=-group['lr'])

return loss

```

这里通过 `grad = p.grad` 得到每个参数的梯度，也就是1式的 g_t 。

如果使用 `weight_decay` 的话，那么相当于目标函数加上 $\frac{1}{2}\lambda||W||^2$ ，所以相当于是梯度相当于要再加上 λW ，故使用了 `grad = grad.add(p, alpha=group['weight_decay'])`。

`square_avg.mul_(alpha).addcmul_(grad, grad, value=1 - alpha)` 对应10式，计算当前步的 V_t 。

`centered` 这一项是 `False` 的话直接 `square_avg.sqrt().add_(group['eps'])` 对 V_t 开根号。

`centered` 这一项是 `True` 的话就把方差使用梯度作归一化。

最后通过 `p.addcdiv_(grad, avg, value=-group['lr'])` 更新梯度，相当于是上面的 3 式。

RMSprop算是Adagrad的一种发展，和Adadelat的变体，效果趋于二者之间

定义优化器：

```
CLASS torch.optim.Adadelta(params, lr=1.0, rho=0.9, eps=1e-06, weight_decay=0)
```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate – 相当于是统一框架中的 α 。
- **rho**(float,optional) – 计算梯度平方的滑动平均超参数 (默认值：0.9)
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty) (默认值：0)
- **eps**(float,optional)：防止分母为0的一个小数 (默认值：1e-10)

源码解读：

```
import torch

from .optimizer import Optimizer

[docs]class Adadelta(Optimizer):
    """Implements Adadelta algorithm.

    It has been proposed in `ADADELTA: An Adaptive Learning Rate Method`__.

    Arguments:
        params (iterable): iterable of parameters to optimize or dicts defining
            parameter groups
        rho (float, optional): coefficient used for computing a running average
            of squared gradients (default: 0.9)
        eps (float, optional): term added to the denominator to improve
            numerical stability (default: 1e-6)
        lr (float, optional): coefficient that scale delta before it is applied
            to the parameters (default: 1.0)
        weight_decay (float, optional): weight decay (L2 penalty) (default: 0)

    __ https://arxiv.org/abs/1212.5701
    """

    def __init__(self, params, lr=1.0, rho=0.9, eps=1e-6, weight_decay=0):
        if not 0.0 <= lr:
```

```

        raise ValueError("Invalid learning rate: {}".format(lr))
    if not 0.0 <= rho <= 1.0:
        raise ValueError("Invalid rho value: {}".format(rho))
    if not 0.0 <= eps:
        raise ValueError("Invalid epsilon value: {}".format(eps))
    if not 0.0 <= weight_decay:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))

    defaults = dict(lr=lr, rho=rho, eps=eps, weight_decay=weight_decay)
    super(Adadelta, self).__init__(params, defaults)

[docs]    @torch.no_grad()
    def step(self, closure=None):
        """Performs a single optimization step.

        Arguments:
            closure (callable, optional): A closure that reevaluates the model
                and returns the loss.
        """
        loss = None
        if closure is not None:
            with torch.enable_grad():
                loss = closure()

        for group in self.param_groups:
            for p in group['params']:
                if p.grad is None:
                    continue
                grad = p.grad
                if grad.is_sparse:
                    raise RuntimeError('Adadelta does not support sparse gradients')
                state = self.state[p]

                # State initialization
                if len(state) == 0:
                    state['step'] = 0
                    state['square_avg'] = torch.zeros_like(p, memory_format=torch.pr
                    state['acc_delta'] = torch.zeros_like(p, memory_format=torch.pr

                square_avg, acc_delta = state['square_avg'], state['acc_delta']
                rho, eps = group['rho'], group['eps']

                state['step'] += 1

                if group['weight_decay'] != 0:
                    grad = grad.add(p, alpha=group['weight_decay'])

                square_avg.mul_(rho).addcmul_(grad, grad, value=1 - rho)
                std = square_avg.add(eps).sqrt_()
                delta = acc_delta.add(eps).sqrt_().div_(std).mul_(grad)
                p.add_(delta, alpha=-group['lr'])
                acc_delta.mul_(rho).addcmul_(delta, delta, value=1 - rho)

        return loss

```

这里通过 `grad = p.grad` 得到每个参数的梯度，也就是1式的 g_t 。

如果使用 `weight_decay` 的话，那么相当于目标函数加上 $\frac{1}{2}\lambda||W||^2$ ，所以相当于梯度相当于要再加上 λW ，故使用了 `grad = grad.add(p, alpha=group['weight_decay'])`。

`square_avg.mul_(rho).addcmul_(grad, grad, value=1 - rho)` 对应10式，计算当前步的 V_t 。`std = square_avg.add(eps).sqrt_()` 对 V_t 开根号。

最后通过 `p.add_(delta, alpha=-group['lr'])` 更新梯度，相当于是上面的 3 式。

`delta` 的分子项是 g_t ，分母项是 V_t 开根号。`acc_delta` 是对 `delta` 的滑动平均。

Adam

谈到这里，Adam和Nadam的出现就很自然而然了——它们是前述方法的集大成者。我们看到，SGD-M在SGD基础上增加了一阶动量，AdaGrad和AdaDelta在SGD基础上增加了二阶动量。把一阶动量和二阶动量都用起来，就是Adam了——Adaptive + Momentum。

SGD的一阶动量：

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

加上AdaDelta的二阶动量：

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad (14)$$

$$\hat{V}_t = \frac{V_t}{1 - \beta_2^t} \quad (15)$$

优化算法里最常见的两个超参数 β_1, β_2 就都在这里了，前者控制一阶动量，后者控制二阶动量。

Nadam

最后是Nadam。我们说Adam是集大成者，但它居然遗漏了Nesterov，这还能忍？必须给它加上，按照NAG的步骤1：

$$g_t = \nabla_{\theta} L(\theta_t - \gamma \hat{m}_{t-1})$$

这就是Nesterov + Adam = Nadam了。

定义优化器：

```
CLASS torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_decay
```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate – 相当于是统一框架中的 α 。
- **betas** (Tuple[float, float], optional) – coefficients used for computing running averages of gradient and its square ((默认值：(0.9, 0.999))
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty) (默认值：0)
- **eps** (float, optional)：防止分母为0的一个小数 (默认值：1e-10)

源码解读：

```
import math
import torch
from .optimizer import Optimizer
```

```
[docs]class Adam(Optimizer):
    r"""Implements Adam algorithm.
```

It has been proposed in `Adam: A Method for Stochastic Optimization`_.

Arguments:

params (iterable): iterable of parameters to optimize or dicts defining parameter groups
lr (float, optional): learning rate (default: 1e-3)
betas (Tuple[float, float], optional): coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))
eps (float, optional): term added to the denominator to improve numerical stability (default: 1e-8)
weight_decay (float, optional): weight decay (L2 penalty) (default: 0)
amsgrad (boolean, optional): whether to use the AMSGrad variant of this algorithm from the paper `On the Convergence of Adam and Beyond`_ (default: False)

```

.. _Adam\: A Method for Stochastic Optimization:
   https://arxiv.org/abs/1412.6980
.. _On the Convergence of Adam and Beyond:
   https://openreview.net/forum?id=ryQu7f-RZ
"""

def __init__(self, params, lr=1e-3, betas=(0.9, 0.999), eps=1e-8,
              weight_decay=0, amsgrad=False):
    if not 0.0 <= lr:
        raise ValueError("Invalid learning rate: {}".format(lr))
    if not 0.0 <= eps:
        raise ValueError("Invalid epsilon value: {}".format(eps))
    if not 0.0 <= betas[0] < 1.0:
        raise ValueError("Invalid beta parameter at index 0: {}".format(betas[0]))
    if not 0.0 <= betas[1] < 1.0:
        raise ValueError("Invalid beta parameter at index 1: {}".format(betas[1]))
    if not 0.0 <= weight_decay:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))
    defaults = dict(lr=lr, betas=betas, eps=eps,
                    weight_decay=weight_decay, amsgrad=amsgrad)
    super(Adam, self).__init__(params, defaults)

def __setstate__(self, state):
    super(Adam, self).__setstate__(state)
    for group in self.param_groups:
        group.setdefault('amsgrad', False)

[docs] @torch.no_grad()
def step(self, closure=None):
    """Performs a single optimization step.

    Arguments:
        closure (callable, optional): A closure that reevaluates the model
            and returns the loss.
    """
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:
                continue
            grad = p.grad
            if grad.is_sparse:
                raise RuntimeError('Adam does not support sparse gradients, please
                    use SparseAdam')

            amsgrad = group['amsgrad']

            state = self.state[p]

            # State initialization
            if len(state) == 0:
                state['step'] = 0

```

```

# Exponential moving average of gradient values
state['exp_avg'] = torch.zeros_like(p, memory_format=torch.preserve_format)
# Exponential moving average of squared gradient values
state['exp_avg_sq'] = torch.zeros_like(p, memory_format=torch.preserve_format)
if amsgrad:
    # Maintains max of all exp. moving avg. of sq. grad. values
    state['max_exp_avg_sq'] = torch.zeros_like(p, memory_format=torch.preserve_format)

exp_avg, exp_avg_sq = state['exp_avg'], state['exp_avg_sq']
if amsgrad:
    max_exp_avg_sq = state['max_exp_avg_sq']
beta1, beta2 = group['betas']

state['step'] += 1
bias_correction1 = 1 - beta1 ** state['step']
bias_correction2 = 1 - beta2 ** state['step']

if group['weight_decay'] != 0:
    grad = grad.add(p, alpha=group['weight_decay'])

# Decay the first and second moment running average coefficient
exp_avg.mul_(beta1).add_(grad, alpha=1 - beta1)
exp_avg_sq.mul_(beta2).addcmul_(grad, grad, value=1 - beta2)
if amsgrad:
    # Maintains the maximum of all 2nd moment running avg. till now
    torch.max(max_exp_avg_sq, exp_avg_sq, out=max_exp_avg_sq)
    # Use the max. for normalizing running avg. of gradient
    denom = (max_exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(1)
else:
    denom = (exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(1)

step_size = group['lr'] / bias_correction1

p.addcdiv_(exp_avg, denom, value=-step_size)

return loss

```

这里通过 $\text{grad} = \text{p.grad}$ 得到每个参数的梯度，也就是1式的 g_t 。

如果使用 `weight_decay` 的话，那么相当于目标函数加上 $\frac{1}{2}\lambda\|W\|^2$ ，所以相当于是梯度相当于要再加上 λW ，故使用了 $\text{grad} = \text{grad.add}(\text{p}, \text{alpha}=\text{group}[\text{'weight_decay'}])$ 。

$\text{exp_avg.mul}(\text{beta1}).\text{add}(\text{grad}, \text{alpha}=1 - \text{beta1})$ 计算12式。

$\text{exp_avg_sq.mul}(\text{beta2}).\text{addcmul}(\text{grad}, \text{grad}, \text{value}=1 - \text{beta2})$ 计算13式。

因为15式的缘故，要给分母除以 $\text{math}^{**}.\text{sqrt}(\text{bias_correction2})$ 。

因为14式的缘故，要给分子除以 `bias_correction1`。

最后通过 $\text{p.addcdiv}(\text{exp_avg}, \text{denom}, \text{value}=-\text{step_size})$ 更新梯度，相当于是上面的 3 式。

AdamW

下图1所示为Adam的另一个改进版：AdamW。

简单来说，AdamW就是Adam优化器加上L2正则，来限制参数值不可太大，这一点属于机器学习入门知识了。以往的L2正则则是直接加在损失函数上，比如这样子：加入正则，损失函数就会变成这样子：

所以在计算梯度 g_t 时要加上粉色的这一项。

但AdamW稍有不同，如下图所示，将正则加在了绿色位置。

Algorithm 2 Adam with L₂ regularization and Adam with decoupled weight decay (AdamW)

1: **given** $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \lambda \in \mathbb{R}$

2: **initialize** time step $t \leftarrow 0$, parameter vector $\theta_{t=0} \in \mathbb{R}^n$, first moment vector $m_{t=0} \leftarrow \mathbf{0}$, second moment vector $v_{t=0} \leftarrow \mathbf{0}$, schedule multiplier $\eta_{t=0} \in \mathbb{R}$

3: **repeat**

4: $t \leftarrow t + 1$

5: $\nabla f_t(\theta_{t-1}) \leftarrow \text{SelectBatch}(\theta_{t-1})$ ▷ select batch and return the corresponding gradient

6: $g_t \leftarrow \nabla f_t(\theta_{t-1}) + \lambda \theta_{t-1}$

7: $m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$ ▷ here and below all operations are element-wise

8: $v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$

9: $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ ▷ β_1 is taken to the power of t

10: $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ ▷ β_2 is taken to the power of t

11: $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$ ▷ can be fixed, decay, or also be used for warm restarts

12: $\theta_t \leftarrow \theta_{t-1} - \eta_t \left(\alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) + \lambda \theta_{t-1} \right)$

13: **until** stopping criterion is met

14: **return** optimized parameters θ_t

图1：AdamW

至于为何这么做？直接摘录BERT里面的原话看看：

Just adding the square of the weights to the loss function is **not** the correct way of using L2 regularization/weight decay with Adam, since that will interact with the m and v parameters in strange ways. Instead we want to decay the weights in a manner that doesn't interact with the m/v parameters. This is equivalent to adding the square of the weights to the loss with plain (non-momentum) SGD. Add weight decay at the end (fixed version).

这段话意思是说，如果直接将L2正则加到loss上去，由于Adam优化器的后序操作，该正则项将会与 m_t 和 v_t 产生奇怪的作用。因而，AdamW选择将 L_2 正则项加在了Adam的 m_t 和 v_t 等参数被计算完之后、在与学习率 η 相乘之前，所以这也表明了weight_decay和 L_2 正则虽目的一致、公式一致，但用法还是不同，二者有着明显的差别。以 PyTorch1.7.0 中的AdamW代码为例：

定义优化器：

```
CLASS torch.optim.AdamW(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_deca
```

参数：

- **params** (iterable) – 优化器作用的模型参数。
- **lr** (float) – learning rate – 相当于是统一框架中的 α 。
- **betas**(Tuple[float,float],optional) – coefficients used for computing running averages of gradient and its square ((默认值：(0.9, 0.999))
- **weight_decay** (float, optional) – 权重衰减系数 weight decay (L2 penalty) (默认值：0)
- **eps**(float,optional)：防止分母为0的一个小数 (默认值：1e-10)

源码解读：

```
import math
import torch
from .optimizer import Optimizer
```

```
[docs]class AdamW(Optimizer):
    r"""Implements AdamW algorithm.
```

```
The original Adam algorithm was proposed in `Adam: A Method for Stochastic Optim
The AdamW variant was proposed in `Decoupled Weight Decay Regularization`_.
```

```
Arguments:
```

```
    params (iterable): iterable of parameters to optimize or dicts defining
        parameter groups
```

```

    lr (float, optional): learning rate (default: 1e-3)
    betas (Tuple[float, float], optional): coefficients used for computing
        running averages of gradient and its square (default: (0.9, 0.999))
    eps (float, optional): term added to the denominator to improve
        numerical stability (default: 1e-8)
    weight_decay (float, optional): weight decay coefficient (default: 1e-2)
    amsgrad (boolean, optional): whether to use the AMSGrad variant of this
        algorithm from the paper `On the Convergence of Adam and Beyond`_
        (default: False)

.. _Adam\: A Method for Stochastic Optimization:
    https://arxiv.org/abs/1412.6980
.. _Decoupled Weight Decay Regularization:
    https://arxiv.org/abs/1711.05101
.. _On the Convergence of Adam and Beyond:
    https://openreview.net/forum?id=ryQu7f-RZ
"""

def __init__(self, params, lr=1e-3, betas=(0.9, 0.999), eps=1e-8,
              weight_decay=1e-2, amsgrad=False):
    if not 0.0 <= lr:
        raise ValueError("Invalid learning rate: {}".format(lr))
    if not 0.0 <= eps:
        raise ValueError("Invalid epsilon value: {}".format(eps))
    if not 0.0 <= betas[0] < 1.0:
        raise ValueError("Invalid beta parameter at index 0: {}".format(betas[0]))
    if not 0.0 <= betas[1] < 1.0:
        raise ValueError("Invalid beta parameter at index 1: {}".format(betas[1]))
    if not 0.0 <= weight_decay:
        raise ValueError("Invalid weight_decay value: {}".format(weight_decay))
    defaults = dict(lr=lr, betas=betas, eps=eps,
                    weight_decay=weight_decay, amsgrad=amsgrad)
    super(AdamW, self).__init__(params, defaults)

def __setstate__(self, state):
    super(AdamW, self).__setstate__(state)
    for group in self.param_groups:
        group.setdefault('amsgrad', False)

[docs] @torch.no_grad()
def step(self, closure=None):
    """Performs a single optimization step.

    Arguments:
        closure (callable, optional): A closure that reevaluates the model
            and returns the loss.
    """
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:

```

```

        continue

    # Perform stepweight decay
    p.mul_(1 - group['lr'] * group['weight_decay'])

    # Perform optimization step
    grad = p.grad
    if grad.is_sparse:
        raise RuntimeError('Adam does not support sparse gradients, please
    amsgrad = group['amsgrad']

    state = self.state[p]

    # State initialization
    if len(state) == 0:
        state['step'] = 0
        # Exponential moving average of gradient values
        state['exp_avg'] = torch.zeros_like(p, memory_format=torch.preserve_format)
        # Exponential moving average of squared gradient values
        state['exp_avg_sq'] = torch.zeros_like(p, memory_format=torch.preserve_format)
        if amsgrad:
            # Maintains max of all exp. moving avg. of sq. grad. values
            state['max_exp_avg_sq'] = torch.zeros_like(p, memory_format=torch.preserve_format)

    exp_avg, exp_avg_sq = state['exp_avg'], state['exp_avg_sq']
    if amsgrad:
        max_exp_avg_sq = state['max_exp_avg_sq']
    beta1, beta2 = group['betas']

    state['step'] += 1
    bias_correction1 = 1 - beta1 ** state['step']
    bias_correction2 = 1 - beta2 ** state['step']

    # Decay the first and second moment running average coefficient
    exp_avg.mul_(beta1).add_(grad, alpha=1 - beta1)
    exp_avg_sq.mul_(beta2).addcmul_(grad, grad, value=1 - beta2)
    if amsgrad:
        # Maintains the maximum of all 2nd moment running avg. till now
        torch.max(max_exp_avg_sq, exp_avg_sq, out=max_exp_avg_sq)
        # Use the max. for normalizing running avg. of gradient
        denom = (max_exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(1)
    else:
        denom = (exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(1)

    step_size = group['lr'] / bias_correction1

    p.addcdiv_(exp_avg, denom, value=-step_size)

    return loss

```

与 Adam 不一样的地方是：

Adam 如果使用 `weight_decay` 的话，那么相当于目标函数加上 $\frac{1}{2}\gamma\|\theta\|^2$ ，所以相

当于是梯度相当于要再加上 $\gamma\theta$ ，故使用了 `grad = grad.add(p, alpha=group['weight_decay'])`。

而 AdamW 是 `p.mul_(1 - group['lr'] * group['weight_decay'])` 直接让参数：

$$\theta_t = \theta_{t-1} - \alpha \cdot \lambda \cdot \theta_{t-1} - \alpha \cdot \eta_t$$

这样才能和绿色框一致

END

我是[王博Kings](#)，985AI博士，华为云专家、CSDN博客专家（人工智能领域优质作者）。单个AI开源项目现在已经获得了2100+标星。现在在做AI相关内容，欢迎一起交流学习、生活各方面的问题，一起加油进步！

我们微信交流群涵盖以下方向（但并不局限于以下内容）：人工智能，计算机视觉，自然语言处理，目标检测，语义分割，自动驾驶，GAN，强化学习，SLAM，人脸检测，最新算法，最新论文，OpenCV，TensorFlow，PyTorch，开源框架，学习方法...

这是我的私人微信，位置有限，一起进步！



王博的公众号，欢迎关注，干货多多



计算机视觉联盟

团队成员为国内各大高校博士，专注于深度学习、机器学习、图像解译、人工...
115篇原创内容

Official Account

王博Kings的系列手推笔记（附高清PDF下载）：

博士笔记 | 周志华《机器学习》手推笔记第一章思维导图

博士笔记 | 周志华《机器学习》手推笔记第二章“模型评估与选择”

博士笔记 | 周志华《机器学习》手推笔记第三章“线性模型”

博士笔记 | 周志华《机器学习》手推笔记第四章“决策树”

博士笔记 | 周志华《机器学习》手推笔记第五章“神经网络”

博士笔记 | 周志华《机器学习》手推笔记第六章支持向量机（上）

博士笔记 | 周志华《机器学习》手推笔记第六章支持向量机（下）